

AI Software Company v1

Your Personal AI Engineering Team

How It Works

- 1 You paste a Problem Statement and click "Start Company"
- 2 CEO creates project brief & assigns work (automatic)
- 3 Business Analyst outputs requirements → You Approve?
- 4 Researcher outputs competitors & APIs → You Approve?
- 5 Architect outputs tech stack & design → You Approve?
- 6 Engineer writes runnable code (only after approvals) → You Approve?
- 7 Presentation Expert creates slides, README, pitch → Done!

Key Features (Built)

- ✓ 6 AI Agents with distinct roles (CEO, BA, Researcher, Architect, Engineer, PPT)
- ✓ Sequential pipeline with 4 founder approval gates
- ✓ Approve or Reject with revision feedback
- ✓ "Call Employee" — direct chat with any agent about the project
- ✓ Shared project memory across all agents
- ✓ Real-time activity feed via WebSockets
- ✓ Visual pipeline progress bar
- ✓ Dark theme dashboard UI
- ✓ OpenRouter integration (use any LLM model)

What Each Agent Produces

CEO	Project name, summary, target users, success criteria, task assignments
Business Analyst	Problem analysis, stakeholders, requirements, user stories, personas, scope, risks
Researcher	Competitors comparison, APIs, open source tools, innovation gaps, recommendation

Architect	System type, tech stack with trade-offs, DB design, API structure, security, scalability
Engineer	Complete runnable code, project structure, setup instructions, dependencies
Presentation	README, pitch deck (8-12 slides), demo flow, technical documentation

Tech Stack

Python + FastAPI	Backend & agent orchestration
Next.js + TypeScript	Frontend dashboard
SQLite	Shared project memory & data
OpenRouter API	LLM access (any model)
WebSockets	Real-time updates
Tailwind CSS	UI styling

Not Yet Built (V2+)

- ✗ Auth / login
- ✗ Multiple projects list
- ✗ Billing / payments
- ✗ Team collaboration
- ✗ Plugin system
- ✗ AI personalities / avatars
- ✗ Code file export / download
- ✗ Streaming agent responses
- ✗ Project history persistence